

---

## Summer 2022 Data Science Intern Challenge

Thi Giang

January 7th, 2022


```
setwd("/Users/gabati/Documents/Internship/Shopify/")
library(readxl)
library(ggplot2)
challenge<-read_excel("Data.xlsx")
```

**Question 1a: Think about what could be going wrong with our calculation. Think about a better way to evaluate this data.**

Answer:

As given in the begining that the shops sell a affordable item—snearker, but the AOV of \$3145.13 per oder is too high. The AOV here was calculated by taking the average revenue of all 100 shops over 5,000 orders (total orders of 100 shops in 30 days). Let have a quick look into the data by observing the boxplot:

```
ggplot(data=data.frame(challenge$order_amount), aes(y=challenge$order_amount))+
  geom_boxplot()+
  scale_x_discrete()+
  coord_trans(y= "log10")+
  ylab("Amount order")+
  scale_y_continuous(breaks=c(100,150,300,400,3000,25000, 70000))
```



There are two things could be going wrong here.

1. First, the mean is extremely sensitive with outliers but this dataset has many potential outliers. There is a shop that sell sneakers with extremely high price (shop 78 with price of \$25,725 per pair). Shop 42 that sales with large quantity per order (2000 items) compared to the majority of the data. It makes these orders become outliers of the dataset. Therefore, if we just naively take the average per order, we will get high AOV but it does not reflect correctly customer buying behavior or their amount purchase each time.

```
table(challenge$total_items)
```

```
##
##      1      2      3      4      5      6      8 2000
## 1830 1832  941  293   77    9    1   17
```

```
cat("Shop that had the most quantity sold in one order is shop", challenge$shop_id[which.max(challenge$
```

```
## Shop that had the most quantity sold in one order is shop 42
```

```
unit_price<-challenge$order_amount/challenge$total_items
```

```
cat("Shop that sold the most expensive item is shop ", challenge$shop_id[which.max(unit_price)], " at $
```

```
## Shop that sold the most expensive item is shop 78 at $25725/ item
```

2. Second, as we all aware that the error might present during collecting the data, I suspect that some observations in this dataset are potentially wrong data. If we look closer on the orders of shop 42, we see that there are multiple orders have the EXACT SAME details about user id, order amount, date and time (see the result below). This may happen when the payment was canceled due to error but the program still recorded the order in database. Plus, these transactions have the highest order amounts of \$704,000 in the whole dataset. So if these transactions are wrong, it obviously indicates that the AOV is misleading.

```
challenge[c(521,4647,61,16,2298,1437,2154,1363,1603,1563,4869,1105,333,4883,2836,2970,4057),]
```

```
## # A tibble: 17 x 7
##   order_id shop_id user_id order_amount total_items payment_method
##   <dbl>   <dbl>   <dbl>       <dbl>       <dbl>   <chr>
## 1     521     42     607       704000       2000 credit_card
## 2    4647     42     607       704000       2000 credit_card
## 3      61     42     607       704000       2000 credit_card
## 4      16     42     607       704000       2000 credit_card
## 5    2298     42     607       704000       2000 credit_card
## 6    1437     42     607       704000       2000 credit_card
## 7    2154     42     607       704000       2000 credit_card
## 8    1363     42     607       704000       2000 credit_card
## 9    1603     42     607       704000       2000 credit_card
## 10   1563     42     607       704000       2000 credit_card
## 11   4869     42     607       704000       2000 credit_card
## 12   1105     42     607       704000       2000 credit_card
## 13    333     57     794         441         3 credit_card
```

```
## 14      4883      42      607      704000      2000 credit_card
## 15      2836      42      607      704000      2000 credit_card
## 16      2970      42      607      704000      2000 credit_card
## 17      4057      42      607      704000      2000 credit_card
## # ... with 1 more variable: created_at <dtm>
```

Think about a better way to evaluate this data.

As I suspected that there might be wrong observations in the data, I would have double checked and cleaned the data.

I think a better way to evaluate this dataset is to separate the outliers and the majority of the data.

**Question 1b** What metric would you report for this dataset?

One quick way to deal with outlier is to use Median which is not affected by outliers. However, if we use mean, it will be skewed by outliers.

Instead, we can separate the data into two groups. The first group include the orders that have regular amount, and the second group include the orders that have outlier amount.

**Question 1c.** What is its value?

If we simply want to choose a better metric, median of the data which is \$284 is a better choice compared to the mean which is \$284.5.

```
cat("Median is $", median(challenge$order_amount), sep="")
```

```
## Median is $284
```

If separate the dataset into groups that has similar characteristic, the AOV for group 1 is \$302.58.

```
shop42_index<-which(challenge$order_amount==704000)
shop78_index<-which(challenge$shop_id==78)
outlier<-c(shop42_index,shop78_index)
group1<-challenge[-outlier,]
cat("The average order value of group 1 is $", round(sum(group1$order_amount)/length(group1$order_amount)), sep="")
```

```
## The average order value of group 1 is $302.58
```

```
AOV_TABLE<-data.frame()
for (i in 1:100) {
  Total_order_amount<-sum(challenge$order_amount[which(challenge$shop_id==i)])
  Shop_id<-i
  Total_order<-length(which(challenge$shop_id==i))
  Total_items<-sum(challenge$total_items[which(challenge$shop_id==i)])
  Price_unit<-Total_order_amount/Total_items
  AOV_by_shop<-round(Total_order_amount/Total_order,2)
  by_shop<-cbind(Shop_id,Total_order,Total_items,Total_order_amount, Price_unit,AOV_by_shop)
  AOV_TABLE<-rbind(AOV_TABLE,by_shop)
}
AOV_TABLE
```

##	Shop_id	Total_order	Total_items	Total_order_amount	Price_unit	AOV_by_shop
## 1	1	44	86	13588	158	308.82
## 2	2	55	102	9588	94	174.33
## 3	3	48	99	14652	148	305.25
## 4	4	51	103	13184	128	258.51
## 5	5	45	92	13064	142	290.31
## 6	6	59	121	22627	187	383.51
## 7	7	56	109	12208	112	218.00
## 8	8	46	84	11088	132	241.04
## 9	9	59	117	13806	118	234.00
## 10	10	53	119	17612	148	332.30
## 11	11	49	95	17480	184	356.73
## 12	12	53	93	18693	201	352.70
## 13	13	63	136	21760	160	345.40
## 14	14	58	121	14036	116	242.00
## 15	15	52	105	16065	153	308.94
## 16	16	41	71	11076	156	270.15
## 17	17	53	100	17600	176	332.08
## 18	18	51	112	17472	156	342.59
## 19	19	64	126	20538	163	320.91
## 20	20	52	103	13081	127	251.56
## 21	21	46	100	14200	142	308.70
## 22	22	48	90	13140	146	273.75
## 23	23	55	112	17472	156	317.67
## 24	24	55	126	17640	140	320.73
## 25	25	48	86	11180	130	232.92
## 26	26	49	95	16720	176	341.22
## 27	27	54	107	18083	169	334.87
## 28	28	43	84	13776	164	320.37
## 29	29	58	118	19234	163	331.62
## 30	30	56	108	16524	153	295.07
## 31	31	47	98	12642	129	268.98
## 32	32	42	79	7979	101	189.98
## 33	33	40	87	15051	173	376.28
## 34	34	50	96	11712	122	234.24
## 35	35	52	104	17056	164	328.00
## 36	36	50	98	12740	130	254.80
## 37	37	48	115	16330	142	340.21
## 38	38	35	72	13680	190	390.86
## 39	39	41	82	10988	134	268.00
## 40	40	48	88	14168	161	295.17
## 41	41	59	127	14986	118	254.00
## 42	42	51	34063	11990176	352	235101.49
## 43	43	58	107	19367	181	333.91
## 44	44	39	71	10224	144	262.15
## 45	45	58	110	15620	142	269.31
## 46	46	43	90	14940	166	347.44
## 47	47	47	84	12180	145	259.15
## 48	48	40	83	9711	117	242.78
## 49	49	53	115	14835	129	279.91
## 50	50	44	92	17756	193	403.55
## 51	51	46	89	16643	187	361.80
## 52	52	41	89	12994	146	316.93
## 53	53	68	130	14560	112	214.12

## 54	54	50	104	13832	133	276.64
## 55	55	48	92	15732	171	327.75
## 56	56	37	69	8073	117	218.19
## 57	57	53	107	15729	147	296.77
## 58	58	59	109	15042	138	254.95
## 59	59	60	121	21538	178	358.97
## 60	60	47	93	16461	177	350.23
## 61	61	50	109	17222	158	344.44
## 62	62	43	83	13280	160	308.84
## 63	63	58	113	15368	136	264.97
## 64	64	43	88	11704	133	272.19
## 65	65	54	116	17864	154	330.81
## 66	66	53	103	16583	161	312.89
## 67	67	37	77	10087	131	272.62
## 68	68	47	88	11968	136	254.64
## 69	69	60	121	15851	131	264.18
## 70	70	59	117	20241	173	343.07
## 71	71	66	130	21320	164	323.03
## 72	72	46	89	14240	160	309.57
## 73	73	58	118	19470	165	335.69
## 74	74	38	76	11628	153	306.00
## 75	75	42	79	10112	128	240.76
## 76	76	42	87	13485	155	321.07
## 77	77	50	90	14040	156	280.80
## 78	78	46	88	2263800	25725	49213.04
## 79	79	54	98	17738	181	328.48
## 80	80	45	93	13485	145	299.67
## 81	81	59	128	22656	177	384.00
## 82	82	42	83	14691	177	349.79
## 83	83	42	81	10449	129	248.79
## 84	84	59	132	20196	153	342.31
## 85	85	35	67	11524	172	329.26
## 86	86	52	111	14430	130	277.50
## 87	87	52	102	15198	149	292.27
## 88	88	50	101	17776	176	355.52
## 89	89	61	118	23128	196	379.15
## 90	90	49	111	19758	178	403.22
## 91	91	54	110	17600	160	325.93
## 92	92	42	76	6840	90	162.86
## 93	93	59	111	12654	114	214.47
## 94	94	45	100	13400	134	297.78
## 95	95	39	74	12432	168	318.77
## 96	96	51	110	16830	153	330.00
## 97	97	48	96	15552	162	324.00
## 98	98	58	107	14231	133	245.36
## 99	99	54	94	18330	195	339.44
## 100	100	40	77	8547	111	213.68

In group 2, the observations are divided into two categories. One is order with high quantity and one is order with high price item. Therefore, I used AOV\_TABLE to extract the AOV of shop 78 which is \$49213.04 per order. This number is more suitable to be the estimation of the average order value for this shop. Shop 42 sells sneaker with price of \$352 per items. This price is on the high side compared to the usual price of sneaker, but it is still affordable. This shop had 51 orders in 30 days window. Among those, 17 orders are big orders with the total sell of 2000 items each. The AOV for this shop is \$235101.49 because of the big

transactions. This shop can actually divide their customer into regular and high spender. They can use the AOV of group 1 to target on customers that spend normally and have some offer, reward program to high spenders to keep them back to the shop.

**Question 2: For this question you'll need to use SQL. Follow this link to access the data set required for the challenge. Please use queries to answer the following questions. Paste your queries along with your final numerical answers below.**

**a. How many orders were shipped by Speedy Express in total?**

Answer: 54 orders were shipped by Speedy Express in total. Solution 1:

```
SELECT COUNT(Orders.OrderID) AS ShipBySpeedy FROM Orders
WHERE Orders.ShipperID=(SELECT Shippers.ShipperID FROM Shippers WHERE ShipperName='Speedy Express');
```

Solution 2 (using join tables)

```
SELECT Count(Orders.ShipperID) AS ShipBySpeedy FROM Orders
JOIN Shippers ON Orders.ShipperID = Shippers.ShipperID
WHERE ShipperName='Speedy Express';
```

**b. What is the last name of the employee with the most orders?**

Answer: The last name of the employee with most orders is Peacock with total of 40 orders.

```
SELECT Employees.LastName
FROM Employees
WHERE Employees.EmployeeID=
  (SELECT EmployeeID
   FROM (SELECT COUNT(Orders.OrderID) AS TotalOrder, Orders.EmployeeID FROM Orders
        GROUP BY Orders.EmployeeID
        ORDER BY TotalOrder DESC
        LIMIT 1));
```

As the question asked for total orders not total quantity so I use Orders table to get the EmployeeID with highest number of order, and get the LastName of that employee from the Employee table.

**c. What product was ordered the most by customers in Germany?**

Answer: Boston Crab Meat is the product which was ordered the most by customers in Germany with 160 orders.

```
SELECT Products.ProductName FROM Products
WHERE Products.ProductID=
  (SELECT ProductID FROM
   (SELECT ProductID, SUM(Quantity) AS TotalQuantity FROM
    ((SELECT Orders.OrderID FROM Orders
      JOIN Customers ON Orders.CustomerID=Customers.CustomerID
      WHERE Customers.Country="Germany") AS GermanyOrders
     JOIN OrderDetails ON GermanyOrders.OrderID=OrderDetails.OrderID)
   GROUP BY ProductID
   ORDER BY TotalQuantity DESC LIMIT 1));
```